

Investigate a Dataset

Udacity Data Analyst Nanodegree - Project 2

Introduction

This notebook investigates the Titanic (<https://www.kaggle.com/c/titanic/data>) dataset containing demographics and passenger information from 891 of the 2224 passengers and crew on board the Titanic.

Questions

The analysis of the Titanic dataset deals mainly with the relationship between survival of an individual and variables such as his:

- sex
- age
- passenger class
- ticket (fare) price
- number of siblings/spouses on board
- number of parents/children on board

Therefore we are investigating the following main question: **Which factors made survival of an individual more likely?**

During the course of analysis we are also looking at the following specific questions:

1. How did sex, age and socio-economic (passenger class / ticket price) status influence survival?
2. How did relationships on board (number of siblings/spouses/parents/children) influence survival?

Resources

- Udacity "Intro to data analysis" material
- Python 3 documentation (<https://docs.python.org/3/>)
- Pandas documentation (<http://pandas.pydata.org/pandas-docs/stable/>)
- NumPy documentation (<http://docs.scipy.org/doc/>)
- Matplotlib documentation (<http://matplotlib.org/contents.html#>)
- Seaborn documentation (<https://stanford.edu/~mwaskom/software/seaborn/>)
- Markdown documentation (<https://daringfireball.net/projects/markdown/syntax>)
- 20 Python libraries you aren't using (but should) (<https://www.oreilly.com/learning/20-python-libraries-you-arent-using-but-should>)

Environment setup

```
In [1]: # Load required modules
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# display plots inside the notebook
%matplotlib inline

# ensure compatibility with Python 2.x
# from __future__ import print_function
```

Data ingestion

```
In [2]: # Load dataset from local file system
titanic = pd.read_csv("titanic_data.csv")
```

Data exploration

Let's explore the dataset by printing its shape, the first and last 5 rows of data, and calculating some summary statistics

```
In [3]: # print shape rows, columns) of data set
titanic.shape
```

```
Out[3]: (891, 12)
```

In [4]: *# show first 5 rows of dataset*
 titanic.head()

Out[4]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22	1	0	A/5 21171	7.25
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38	1	0	PC 17599	71.28
2	3	1	3	Heikkinen, Miss. Laina	female	26	0	0	STON/O2. 3101282	7.92
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35	1	0	113803	53.10
4	5	0	3	Allen, Mr. William Henry	male	35	0	0	373450	8.05



In [5]: *# show last 5 rows of dataset*
 titanic.tail()

Out[5]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
886	887	0	2	Montvila, Rev. Juozas	male	27	0	0	211536	13.00
887	888	1	1	Graham, Miss. Margaret Edith	female	19	0	0	112053	30.00
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.45
889	890	1	1	Behr, Mr. Karl Howell	male	26	0	0	111369	30.00
890	891	0	3	Dooley, Mr. Patrick	male	32	0	0	370376	7.75



Summary statistics

Looking at each variable indepdently the summary statistics tell us that:

- PassengerId: There were 891 passengers on board
- Survived: Only 38% of these passengers survived.
- PClass: Only few passengers could afford first class, most, about 50%, spent their time on board in third class
- Age: Only 714 observations contain information for age, could this variable be a potential candidate for a data cleaning exercise? Mean age 30 with high standard deviation, youngest passenger still a baby, oldest passenger an old person at age 80. He or she seems to be an extreme outlier, since the average age in the third percentile is 38.
- SibSp: Mean value of 0.52 is somewhat misleading, since one can't have half a sibling or family member on board. Interestingly the standard deviation is quite high. Looking at the max value, there seems to be either a large family on board or someone married to many wives (max = 8)
- Parch: There seems to be a large family onboard (max = 6)
- Fare: Mean price for a titanic ticket was 32 USD, although the standard deviationn is quite high (USD 50). Apparently some passengers did not pay anything for their ticket (min = 0), while some potentially wealthy passengers paid up to 512 USD (max = 512).

Given the summary statistics we might investigate the following questions:

- What does high standard deviation of Survied mean? Are their differences between gender, age, socio-economic status, etc.?
- Who were the youngest and oldest passenger on board (Age = 0.42/80)
- Who belonged to the large family?
- Who paid nothing at all and the maxium price?

```
In [6]: # calculate summary statistics
titanic.describe()
```

```
Out[6]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.200726
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910461
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454243
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.001754
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.3292

Helper functions

```
In [7]: # helper functions to print rows containing the min/max value of a variable
def titanic_min(variable):
    """
    Given a variable present in the titanic data set, the function prints the
    rows containing the min value
    """
    print("Information for min values of %s:" % variable)
    print(titanic.ix[titanic[variable] == min(titanic[variable])])

def titanic_max(variable):
    """
    Given a variable present in the titanic data set, the function prints the
    row containing the max value
    """
    print("Information for max values of %s:" % variable)
    print(titanic.ix[titanic[variable] == max(titanic[variable])])
```

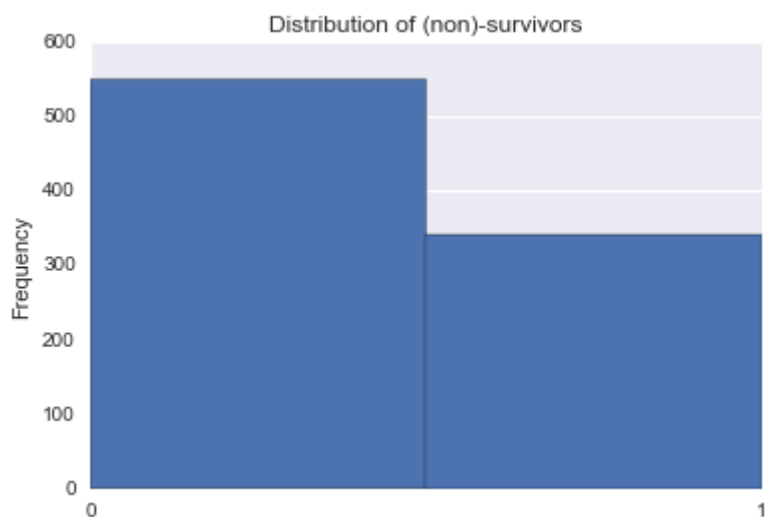
Visualization

Besides looking at plain figures, we also like to investigate our data visually.

Let's look at **survival** first. From the bar chart below we can see that only about 350 of 891 passengers survived their trip.

```
In [8]: # plot survival data
titanic["Survived"].plot(kind="hist", title="Distribution of (non)-survivors",
    bins=2, xticks=(0,1))
```

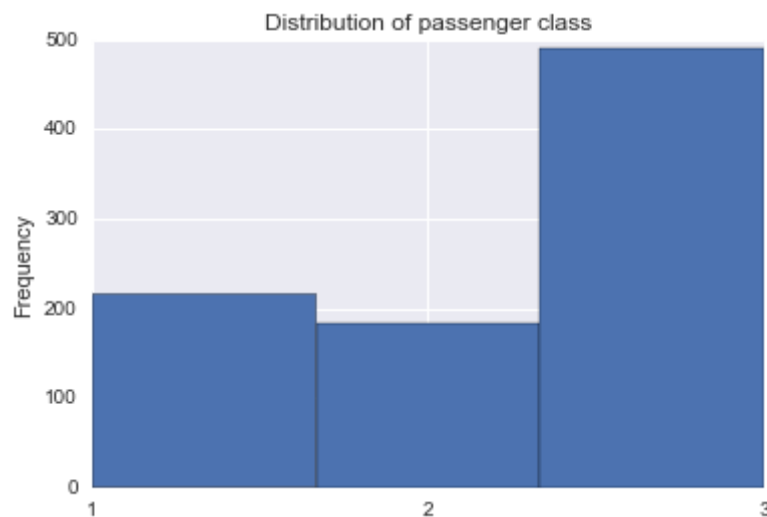
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0xacefc50>



What about **passenger class**? Apparently half of passengers were traveling in third class. The other half almost equally split into second and first class.

```
In [9]: # plot passenger class data
titanic["Pclass"].plot(kind="hist", title="Distribution of passenger class", bins=3, xticks=(1,2,3))
```


```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0xb0c0dd8>
```



Next, let's investigate the **age** distribution of Titanic passengers. Apparently most of the passengers were between 20 and 30 years old. From the histogram it is evident that there were some very old passengers, too.

```
In [10]: # plot age data
titanic["Age"].plot(kind="hist", title="Distribution of age")
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0xb594550>
```



Now let's have a look at extreme ages. How old were the youngest/oldest passengers? As can be seen from the output below, the youngest passenger was **not even 1 year old**, while the oldest passenger was already **80**. Interestingly both survived, despite travelling in different passenger classes.

```
In [11]: # print information about youngest passenger
titanic_min("Age")
```

Information for min values of Age:

PassengerId	Survived	Pclass	Name	Sex	\
803	804	1	3	Thomas, Master. Assad Alexander	male

Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	
803	0.42	0	1	2625	8.5167	NaN	C

```
In [12]: # print information about oldest passenger
titanic_max("Age")
```

Information for max values of Age:

PassengerId	Survived	Pclass	Name	\
630	631	1	1	Barkworth, Mr. Algernon Henry Wilson

Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	
630	male	80	0	0	27042	30	A23	S

Let's look at distribution of siblings/spouses of Titanic passengers. Interestingly, most passengers either did not have any or just one siblings/spouses on board, while there was one family (or someone with a lot of spouses) with **8** relatives on board.

```
In [13]: # plot sibling/spouse data
titanic["SibSp"].plot(kind="hist", title="Distribution of siblings/spouses on board")
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0xb5edef0>
```



Now who was the family with maximum number of relatives?. As can be seen from the table below, it was the **Sage** family, which unfortunately did not survive their journey.

```
In [14]: # print information about max/min values
         titanic_max("SibSp")
```

Information for max values of SibSp:


	PassengerId	Survived	Pclass	Name	Sex
159	160	0	3	Sage, Master. Thomas Henry	male
180	181	0	3	Sage, Miss. Constance Gladys	female
201	202	0	3	Sage, Mr. Frederick	male
324	325	0	3	Sage, Mr. George John Jr	male
792	793	0	3	Sage, Miss. Stella Anna	female
846	847	0	3	Sage, Mr. Douglas Bullen	male
863	864	0	3	Sage, Miss. Dorothy Edith "Dolly"	female

	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
159	NaN	8	2	CA. 2343	69.55	NaN	S
180	NaN	8	2	CA. 2343	69.55	NaN	S
201	NaN	8	2	CA. 2343	69.55	NaN	S
324	NaN	8	2	CA. 2343	69.55	NaN	S
792	NaN	8	2	CA. 2343	69.55	NaN	S
846	NaN	8	2	CA. 2343	69.55	NaN	S
863	NaN	8	2	CA. 2343	69.55	NaN	S

What about the distribution of parents/children onboard of Titanic? The figure below shows that the majority of passengers did not have any children on board. As seen within the siblings/spouses data, there is one extreme case which we investigate below.

```
In [15]: # plot parent/child data
titanic["Parch"].plot(kind="hist", title="Distribution of parents/children on board")
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0xb68f978>
```



Apparently, Mrs. Goodwin was accompanied by **6 children** and unfortunately did not survive her trip.

```
In [16]: # print information about min/max parent/child data
titanic_max("Parch")
```

Information for max values of Parch:

	PassengerId	Survived	Pclass	Name
\	678	0	3	Goodwin, Mrs. Frederick (Augusta Tyler)

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
678	female	43	1	6	CA 2144	46.9	NaN	S

Finally, let's dig into the distribution of fare prices. Obviously most passengers paid well below USD 100 for their ticket. There are a some passengers who paid more, e.g. between USD 100 and USD 300 while a few payed as much as USD 500.

```
In [17]: # plot fare data  
titanic["Fare"].plot(kind="hist", title="Distribution of fare prices")
```


```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0xb719e80>
```



From the boxplot below we can see that the median fare price was well below USD 100 (from cell 5 we actually know that the average is USD 32 with standard deviation of approximately USD 50. Furthermore the fare price of roughly USD 500 seems to be an outlier.

```
In [18]: # plot fare data as box plot  
ax = sns.boxplot(titanic["Fare"], orient="h")  
ax.set_title("Fare prices")
```

```
Out[18]: <matplotlib.text.Text at 0xb83f320>
```



Now let's check the minimum fare price. Interestingly the minimum price is USD 0, meaning that 15 passengers did not pay for their ticket at all.

```
In [19]: # print passengers who paid the minimum fare price
titanic_min("Fare")
```

Information for min values of Fare:

	PassengerId	Survived	Pclass	Name	Sex	\
179	180	0	3	Leonard, Mr. Lionel	male	
263	264	0	1	Harrison, Mr. William	male	
271	272	1	3	Tornquist, Mr. William Henry	male	
277	278	0	2	Parkes, Mr. Francis "Frank"	male	
302	303	0	3	Johnson, Mr. William Cahoon Jr	male	
413	414	0	2	Cunningham, Mr. Alfred Fleming	male	
466	467	0	2	Campbell, Mr. William	male	
481	482	0	2	Frost, Mr. Anthony Wood "Archie"	male	
597	598	0	3	Johnson, Mr. Alfred	male	
633	634	0	1	Parr, Mr. William Henry Marsh	male	
674	675	0	2	Watson, Mr. Ennis Hastings	male	
732	733	0	2	Knight, Mr. Robert J	male	
806	807	0	1	Andrews, Mr. Thomas Jr	male	
815	816	0	1	Fry, Mr. Richard	male	
822	823	0	1	Reuchlin, Jonkheer. John George	male	

	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
179	36	0	0	LINE	0	NaN	S
263	40	0	0	112059	0	B94	S
271	25	0	0	LINE	0	NaN	S
277	NaN	0	0	239853	0	NaN	S
302	19	0	0	LINE	0	NaN	S
413	NaN	0	0	239853	0	NaN	S
466	NaN	0	0	239853	0	NaN	S
481	NaN	0	0	239854	0	NaN	S
597	49	0	0	LINE	0	NaN	S
633	NaN	0	0	112052	0	NaN	S
674	NaN	0	0	239856	0	NaN	S
732	NaN	0	0	239855	0	NaN	S
806	39	0	0	112050	0	A36	S
815	NaN	0	0	112058	0	B102	S
822	38	0	0	19972	0	NaN	S

```
In [20]: # print number of passengers with minimum ticket price
len(titanic[titanic["Fare"] == 0])
```

Out[20]: 15

What about the maximum fare price? Obviously three passengers were willing to pay the maximum price of **USD 512**, which is 16 times higher than the average price USD 32. At least all three got a ticket for the first passenger class!

```
In [21]: # print passengers who paid the maximum fare price
titanic_max("Fare")
```

Information for max values of Fare:

PassengerId	Survived	Pclass	Name \	
258	259	1	1	Ward, Miss. Anna
679	680	1	1	Cardeza, Mr. Thomas Drake Martinez
737	738	1	1	Lesurer, Mr. Gustave J

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
258	female	35	0	0	PC 17755	512.3292	NaN	C
679	male	36	0	1	PC 17755	512.3292	B51 B53 B55	C
737	male	35	0	0	PC 17755	512.3292	B101	C

Data cleaning

Before moving on to actual analysis the data needs to be cleaned. During the exploration phase we discovered missing values for **Age** and **Cabin**. Furthermore some passengers were not assigned a proper ticket ID, but the value "Line". Another candidate for cleaning could be various extreme values in fare price, siblings/spouses or parents/children. How do we decide which values to keep and which to clean? One approach would be to go back to our initial question and check whether missing values in particular columns could impede analysis. As we are primary interested in factors influencing **survival**, e.g. sex, age, passenger class and other socio-economic variables, we should focus on these during data cleaning

Let's start with investigating real missing values: Age information is missing for 20% of all passengers, while cabin information is missing for **77%** of all passengers. Why do we have so little information on cabins?

```
In [22]: # for each column print number of records where information is missing
titanic.isnull().sum()
```

```
Out[22]: PassengerId      0
Survived      0
Pclass        0
Name          0
Sex           0
Age          177
SibSp         0
Parch         0
Ticket        0
Fare          0
Cabin        687
Embarked      2
dtype: int64
```

```
In [23]: # for each column print missing values as percentage of total values
titanic.isnull().sum() / titanic.shape[0]
```

```
Out[23]: PassengerId    0.000000
Survived              0.000000
Pclass               0.000000
Name                 0.000000
Sex                  0.000000
Age                  0.198653
SibSp                0.000000
Parch                0.000000
Ticket              0.000000
Fare                 0.000000
Cabin                0.771044
Embarked             0.002245
dtype: float64
```

Let's dig deeper into missing age and cabin data. Checking passengers travelling in third class for missing data reveals that most of our issues can be found there. **77%** of missing age and **70%** of missing cabin values are attached to passengers in the third class.

```
In [24]: # print a subset of records with missing age information
titanic.ix[titanic["Age"].isnull()].head()
```

```
Out[24]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
5	6	0	3	Moran, Mr. James	male	NaN	0	0	330877	8.45
17	18	1	2	Williams, Mr. Charles Eugene	male	NaN	0	0	244373	13.00
19	20	1	3	Masselmani, Mrs. Fatima	female	NaN	0	0	2649	7.22
26	27	0	3	Emir, Mr. Farred Chehab	male	NaN	0	0	2631	7.22
28	29	1	3	O'Dwyer, Miss. Ellen "Nellie"	female	NaN	0	0	330959	7.87

```
In [25]: # print a subset of records with missing cabin information
titanic.ix[titanic["Cabin"].isnull()].head()
```

Out[25]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22	1	0	A/5 21171	7.25
2	3	1	3	Heikkinen, Miss. Laina	female	26	0	0	STON/O2. 3101282	7.92
4	5	0	3	Allen, Mr. William Henry	male	35	0	0	373450	8.05
5	6	0	3	Moran, Mr. James	male	NaN	0	0	330877	8.45
7	8	0	3	Palsson, Master. Gosta Leonard	male	2	3	1	349909	21.0

```
In [26]: # for each column where PClass is equal to 3, print number of records where in
formation is missing
titanic[titanic["Pclass"] == 3].isnull().sum()
```

```
Out[26]: PassengerId      0
Survived      0
Pclass        0
Name          0
Sex           0
Age          136
SibSp         0
Parch         0
Ticket        0
Fare          0
Cabin        479
Embarked      0
dtype: int64
```

```
In [27]: # for each column where Pclass is equal to 3, print missing values as percent
age of total values
titanic[titanic["Pclass"] == 3].isnull().sum() / titanic.isnull().sum()
```

```
Out[27]: PassengerId      NaN
Survived      NaN
Pclass        NaN
Name          NaN
Sex           NaN
Age           0.768362
SibSp         NaN
Parch         NaN
Ticket        NaN
Fare          NaN
Cabin         0.697234
Embarked      0.000000
dtype: float64
```

What could be a possible explanation for that? Apparently third class had bunk beds for 4-6 people. Maybe data was not rigorously recorded for this class, see: <https://nmni.com/titanic/On-Board/Sleeping.aspx> (<https://nmni.com/titanic/On-Board/Sleeping.aspx>)

Although there does not seem to be a substantial problem with Embarked and Ticket information, let's have a brief look at the missing values:

```
In [28]: # print records with missing embarked information
titanic.ix[titanic["Embarked"].isnull()]
```

```
Out[28]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	C
61	62	1	1	Icard, Miss. Amelie	female	38	0	0	113572	80	B
829	830	1	1	Stone, Mrs. George Nelson (Martha Evelyn)	female	62	0	0	113572	80	B


```
In [29]: # print records where the value for ticket is "line"
titanic.ix[titanic["Ticket"] == "LINE"]
```

Out[29]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
179	180	0	3	Leonard, Mr. Lionel	male	36	0	0	LINE	0	Na
271	272	1	3	Tornquist, Mr. William Henry	male	25	0	0	LINE	0	Na
302	303	0	3	Johnson, Mr. William Cahoone Jr	male	19	0	0	LINE	0	Na
597	598	0	3	Johnson, Mr. Alfred	male	49	0	0	LINE	0	Na

Now, which data we want to omit for the analysis? Since our analysis mainly focuses on personal data like age, and socio-economic information, we first **remove all columns containing data which does not help to investigate these variables**, such as:

- PassengerId
- Name
- Ticket
- Cabin
- Embarked

```
In [30]: # remove PassengerId, Name, Ticket, Cabin and Embarked column
titanic_cleaned = titanic.drop(labels=["PassengerId", "Name", "Ticket", "Cabin", "Embarked"], axis=1, inplace=False)
```

Further we **remove the outlier values (max values) for fare prices**:

```
In [31]: # remove fare price outliers
titanic_cleaned.drop(titanic_cleaned.index[[258, 679, 737]], inplace=True)

# verify that used-to-be max values for Fare have been removed
titanic_cleaned["Fare"].max()
```

Out[31]: 263.0

We do also **drop all rows not containing age information**:

```
In [32]: # remove any records with missing age information
titanic_cleaned.dropna(subset=["Age"], inplace=True)
```

Finally we check if any missing values and remain, which is not the case. We are left with **711 rows** of cleaned data:

```
In [33]: # check whether rows with missing age information were succesfully removed
titanic_cleaned.isnull().sum()
```

```
Out[33]: Survived    0
Pclass      0
Sex         0
Age         0
SibSp       0
Parch       0
Fare        0
dtype: int64
```

```
In [34]: # print number of rows remaining after cleaning
titanic_cleaned.shape[0]
```

```
Out[34]: 711
```

Analysis

After exploring and cleaning the data, we are finally able to analyze our main question: **Which factors made survival of an individual more likely?** In order to start investigating this question, we would like to know how **survival** is correlated with other variables in the dataset. Although this does not imply causation, i.e. a strong positive correlation between **survival** and travelling in **the first passenger class** does **not proof**, that passengers survived because they travelled in first class. Maybe first class passengers where particulary wealthy and could afford personal that saved them in case of emergency. Taking this into account, using correlation between variables is still a good start for deeper analysis.

```
In [234]: titanic_cleaned.head()
```

```
Out[234]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare
0	0	3	male	22.0	1	0	7.2500
1	1	1	female	38.0	1	0	71.2833
2	1	3	female	26.0	0	0	7.9250
3	1	1	female	35.0	1	0	53.1000
4	0	3	male	35.0	0	0	8.0500

Let's start with age. Instead of correlating Survived with individual ages, we form three age groups (young, middle and old) and use these for analysis:

In [35]: `# bin age into young, middle, old buckets`
`titanic_cleaned["age"] = pd.cut(titanic["Age"], bins=3, labels=["young", "middle", "old"])`
`titanic_cleaned.head()`

Out[35]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	age	fare
0	0	3	male	22	1	0	7.2500	young	
1	1	1	female	38	1	0	71.2833	middle	
2	1	3	female	26	0	0	7.9250	young	
3	1	1	female	35	1	0	53.1000	middle	
4	0	3	male	35	0	0	8.0500	middle	

We apply the same logic to fare prices:

In [37]: `# bin fare prices into low, medium, high buckets`
`titanic_cleaned["fare"] = pd.cut(titanic["Fare"], bins=3, labels=["low", "medium", "high"])`
`titanic_cleaned.head()`

Out[37]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	age	fare
0	0	3	male	22	1	0	7.2500	young	low
1	1	1	female	38	1	0	71.2833	middle	low
2	1	3	female	26	0	0	7.9250	young	low
3	1	1	female	35	1	0	53.1000	middle	low
4	0	3	male	35	0	0	8.0500	middle	low

Further, instead of correlating individual values for siblings/spouses and parents/children, we create the dummy variable "family" to indicate whether a passenger had at least 1 sibling/spouse or 1 parent/child on board.

In [38]: `# create dummy variable "family" to indicate whether a passenger had at least one sibling/spouse OR parent/child on board`
`titanic_cleaned["family"] = (titanic["SibSp"] >= 1) | (titanic["Parch"] >= 1)`
`titanic_cleaned.head()`

Out[38]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	age	fare	family
0	0	3	male	22	1	0	7.2500	young	low	True
1	1	1	female	38	1	0	71.2833	middle	low	True
2	1	3	female	26	0	0	7.9250	young	low	False
3	1	1	female	35	1	0	53.1000	middle	low	True
4	0	3	male	35	0	0	8.0500	middle	low	False

Finally, we convert our recently created variables (age, fare, family), as well as Sex and passenger class into dummy variables.

In [39]: `# convert categorical variables (Sex and PClass) into dummy variables for analysis`
`titanic_dummies = pd.get_dummies(data=titanic_cleaned, columns=["Sex", "Pclass", "age", "family", "fare"])`
`titanic_dummies.drop(labels=["Age", "SibSp", "Parch", "Fare"], axis=1, inplace=True)`
`titanic_dummies.head()`

Out[39]:

	Survived	Sex_female	Sex_male	Pclass_1	Pclass_2	Pclass_3	age_young	age_mid
0	0	0	1	0	0	1	1	0
1	1	1	0	1	0	0	0	1
2	1	1	0	0	0	1	1	0
3	1	1	0	1	0	0	0	1
4	0	0	1	0	0	1	0	1

Now we are ready to calculate the correlation matrix.

```
In [248]: # calculate correlation matrix
titanic_dummies.corr()
```


Out[248]:

	Survived	Sex_female	Sex_male	Pclass_1	Pclass_2	Pclass_3	age_young
Survived	1.000000	0.542190	-0.542190	0.294366	0.085113	-0.330442	0.012834
Sex_female	0.542190	1.000000	-1.000000	0.116341	0.068353	-0.160388	0.064173
Sex_male	-0.542190	-1.000000	1.000000	-0.116341	-0.068353	0.160388	-0.064173
Pclass_1	0.294366	0.116341	-0.116341	1.000000	-0.333844	-0.587952	-0.244611
Pclass_2	0.085113	0.068353	-0.068353	-0.333844	1.000000	-0.566204	-0.067082
Pclass_3	-0.330442	-0.160388	0.160388	-0.587952	-0.566204	1.000000	0.271461
age_young	0.012834	0.064173	-0.064173	-0.244611	-0.067082	0.271461	1.000000
age_middle	0.011208	-0.032636	0.032636	0.126100	0.067855	-0.168494	-0.867727
age_old	-0.046742	-0.060886	0.060886	0.228772	-0.002110	-0.198235	-0.248710
family_False	-0.201643	-0.283819	0.283819	-0.132074	-0.030312	0.141502	-0.113606
family_True	0.201643	0.283819	-0.283819	0.132074	0.030312	-0.141502	0.113606
fare_low	-0.099115	-0.112146	0.112146	-0.250056	0.083480	0.147021	-0.025115
fare_medium	0.099115	0.112146	-0.112146	0.250056	-0.083480	-0.147021	0.025115
fare_high	NaN	NaN	NaN	NaN	NaN	NaN	NaN

Since the correlation matrix is hard to read, we visualize its result using a heatmap.

```
In [40]: # visualize correlation matrix
sns.heatmap(titanic_dummies.corr())
```

```
Out[40]: <matplotlib.axes._subplots.AxesSubplot at 0xb8ac4e0>
```



Interestingly, we can observe a strong (≥ 0.4) positive correlation between **survival** and being female, whereas the opposite is true for being **male**. Further, the correlation matrix shows a moderate positive relation between survival and travelling in **first class**. The opposite is true for residing in **third class**. As far as **age groups** are concerned, no correlation can be observed, while having at least one **family** member on board is modestly correlated with survival. Finally, **fare price groups** do not seem to have a particular strong positive or negative with survival.

Given the positive relation between survival and being female, between survival and passenger class, as well as between survival and family, we look at these variables using a pivot tables.

Solely comparing survival between sex reveals that women were much more likely to survive than men. Further, passengers travelling in first class were much more likely than those travelling in third class. Finally, half of passengers having a family member on board survived their trip, whereas only 1/3 survived without family support.

In [42]: `# pivot table displaying survival vs. sex`
`pd.pivot_table(titanic_cleaned, values=["Survived"], index=["Sex"], aggfunc=[np.sum, np.mean, np.std], margins=True)`

Out[42]:

	sum	mean	std
	Survived	Survived	Survived
Sex			
female	196	0.753846	0.43160
male	91	0.201774	0.40177
All	287	0.403657	0.49063

In [43]: `# pivot table displaying survival vs. passenger class`
`pd.pivot_table(titanic_cleaned, values=["Survived"], index=["Pclass"], aggfunc=[np.sum, np.mean, np.std], margins=True)`

Out[43]:

	sum	mean	std
	Survived	Survived	Survived
Pclass			
1	119	0.650273	0.478192
2	83	0.479769	0.501041
3	85	0.239437	0.427342
All	287	0.403657	0.490630

In [44]: `# pivot table displaying survival vs. family`
`pd.pivot_table(titanic_cleaned, values=["Survived"], index=["family"], aggfunc=[np.sum, np.mean, np.std], margins=True)`

Out[44]:

	sum	mean	std
	Survived	Survived	Survived
family			
False	128	0.318408	0.466439
True	159	0.514563	0.500599
All	287	0.403657	0.490630

Let's go one step further and investigate survival, sex and passenger class. The table shows that **96%** of **females** travelling in first, and **92%** travelling in **second** class survived their trip. Whereas only **40%** of males in first, and only **15%** in second and third class survived their trip.

In [211]: `# pivot table displaying survival vs. sex and passenger class`
`pd.pivot_table(titanic_cleaned, values=["Survived"], index=["Sex", "Pclass"],`
`aggfunc=[np.sum, np.mean, np.std], margins=True)`

Out[211]:

		sum	mean	std
		Survived	Survived	Survived
Sex	Pclass			
female	1	82.0	0.964706	0.185617
	2	68.0	0.918919	0.274823
	3	47.0	0.460784	0.500921
male	1	40.0	0.396040	0.491512
	2	15.0	0.151515	0.360375
	3	38.0	0.150198	0.357973
All		290.0	0.406162	0.491116

Repeating the same analysis using family and passenger class, we can observe relatively high survival rates for passengers travelling in first class, despite their family status. Although, the difference between passengers having a family member on board and travelling in first class, and those without family support is **10%**. However, the major differences here is between passengers having family and travelling in **second class**. While **63%** of passengers having a family member on board and travelling second class survived the trip, only **34%** without family support did.

In [41]: `# pivot table displaying survival vs. family`
`pd.pivot_table(titanic_cleaned, values=["Survived"], index=["family",`
`"Pclass"], aggfunc=[np.sum, np.mean, np.std], margins=True)`

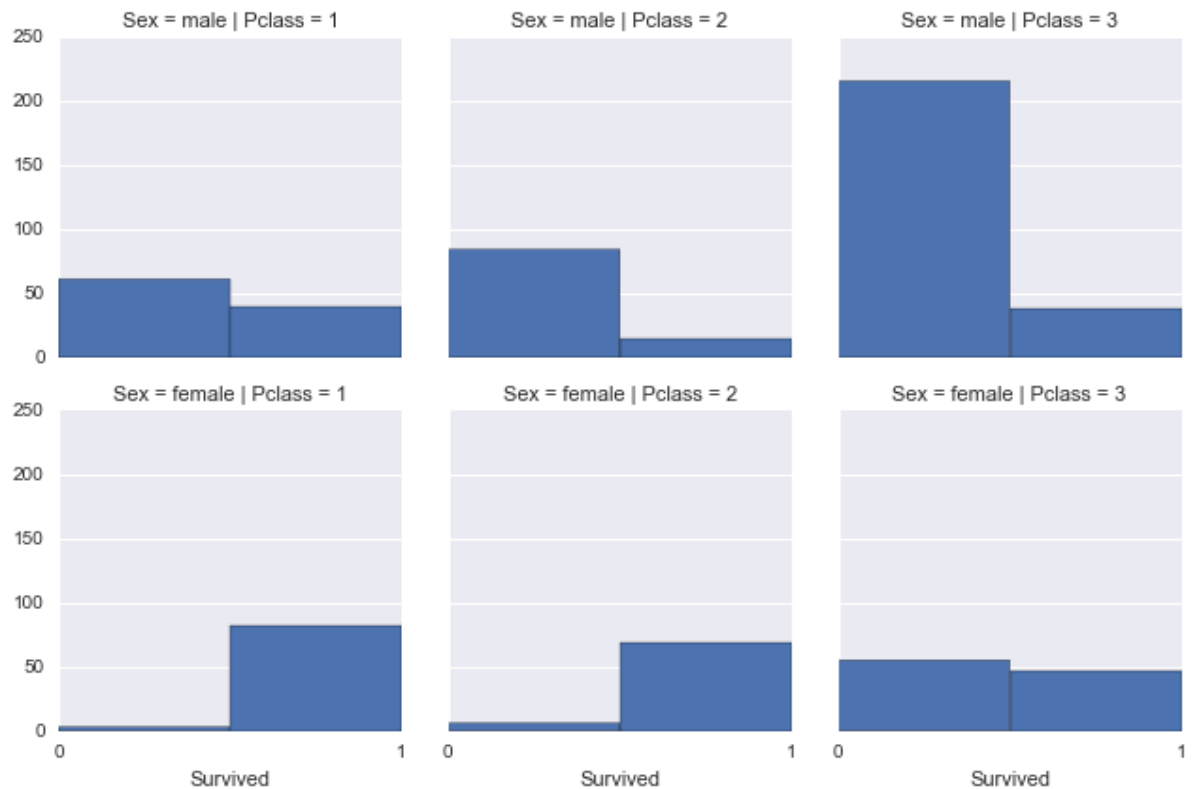
Out[41]:

		sum	mean	std
		Survived	Survived	Survived
family	Pclass			
False	1	49	0.583333	0.495968
	2	32	0.344086	0.477644
	3	47	0.208889	0.407421
True	1	70	0.707071	0.457422
	2	51	0.637500	0.483755
	3	38	0.292308	0.456582
All		287	0.403657	0.490630

In order to make it easier to consume our results, we visualize the analysis above:

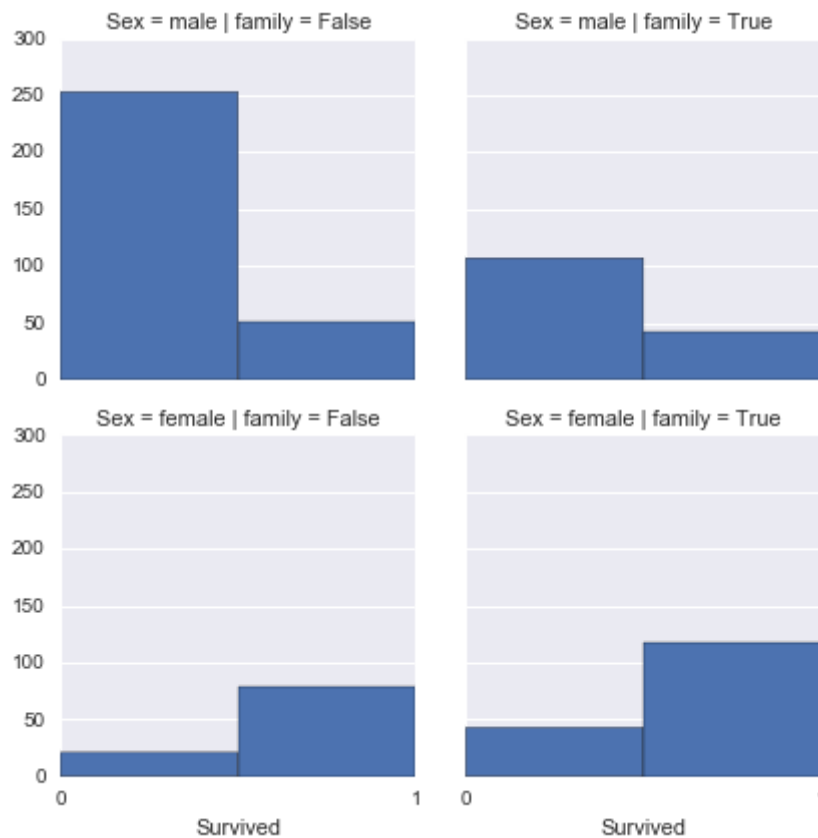

```
In [116]: # visualize differences between (non)-survivors given sex and passenger class  
grid = sns.FacetGrid(titanic_cleaned, row="Sex", col="Pclass")  
grid.map(plt.hist, "Survived", bins=2).set(xticks=(0,1))
```

Out[116]: <seaborn.axisgrid.FacetGrid at 0x1254e2320>



```
In [203]: # visualize differences between (non)-survivors given sex and family on board
grid = sns.FacetGrid(titanic_cleaned, row="Sex", col="family")
grid.map(plt.hist, "Survived", bins=2).set(xticks=(0,1))
```

```
Out[203]: <seaborn.axisgrid.FacetGrid at 0x126218b70>
```



Conclusion

Our general analysis revealed the following points:

- On average women were much more likely to survive their trip on Titanic
- The same is true for passengers travelling in first class
- Passengers having one or more family member on board also survived their trip more often than those without family support

Further, the investigation between survival and more than one variable revealed that:

- Most of the women traveling in first and second class survived their trip, whereas the opposite is true for men
- The data on survival, sex and family on board is inconclusive, since women were on average more likely to survive their trip despite having family on board

Despite our findings, we can not be sure that the variables we found to have an effect on survival, really **caused** it. For instance, there may have been other reasons despite gender, for women to be more likely to survive than men.